James Ding

604-842-6134 | james@dingjames.com| LinkedIn| Github| Portfolio
Calgary, AB

WORK EXPERIENCE

Associate Data Analyst **May 2025 – Present**
AirSprint Inc. Calgary, AB

- Conducted EDA on 12 years of flight operations data (78K+ flights), identified demand drivers and a COVID-era regime shift, then built and deployed a Prophet forecasting model achieving 2% monthly MAPE in production — targeting 1%+ reduction in $49M annual positioning costs
- Built automated data ingestion pipelines (FL3XX, Salesforce, Foreflight → AWS Lambda → S3 → Snowflake) with end-to-end data quality monitoring and alerting
- Designed a data-driven RBAC workaround for FL3XX's booking platform by building a cleansing and segmentation pipeline with custom business logic mapping user roles (owner, individual, full access) and a manual override layer to manage edge cases
- Built QuickSight dashboards enabling real-time monitoring of operational KPIs, flight demand trends, and fleet utilization for executive and operational stakeholders
- Designed and built Jetty, a multi-agent RAG AI assistant for 400+ employees using AWS Lambda, Bedrock, and DynamoDB — enabling natural-language Q&A over internal documentation, policies, and operational data with tool-use and agent orchestration
- Architected an end-to-end vendor contract management pipeline for Legal — automated document intake, processing, and structured storage via custom Salesforce objects, data quality validation, and governed data models replacing fully manual workflows

EDUCATION

Master of Data Science & Analytics **Dec 2025**
University of Calgary Calgary, AB

Bachelor of Science, Agriculture Commerce **May 2016**
University of Illinois at Urbana Champaign Champaign, IL

TECHNICAL SKILLS

Languages: Python, SQL, JavaScript, Go, R
Machine Learning: XGBoost, Random Forest, scikit-learn, time-series forecasting, classification, feature engineering
Deep Learning: PyTorch, TensorFlow, Hugging Face, CNNs, LSTMs, Reinforcement Learning (PPO)
Data Engineering: Data modeling, dimensional modeling, ETL/ELT pipelines, data quality validation, data governance, Snowflake, DuckDB
Platforms: AWS (Lambda, S3, Bedrock, SageMaker, QuickSight, DynamoDB, API Gateway, CDK), Snowflake, Salesforce, Power BI, Git, GitHub Actions, Docker
AI/LLMs: AI agent development (RAG, multi-agent orchestration, tool use), LLM application design (prompt engineering, agent pipelines), AWS Bedrock, OpenAI-compatible APIs